

## ***Kubernetes Ingress***

Ingress is used to expose HTTP and HTTPS routes to external traffic from outside the Kubernetes cluster. It supports the following key features:

- **Path-based routing:** Routes traffic to specific services based on the URL path (e.g., `example.com/app1`).
- **Host-based routing:** Routes traffic based on the hostname or subdomain (e.g., `app.example.com`).
- **Load balancing:** Distributes traffic among multiple pods of a service.
- **SSL termination:** Provides secure communication by terminating SSL/TLS.

Ingress redirects incoming requests to the appropriate services within the cluster based on the web URL or path. It also provides encryption features and helps balance application load effectively.

### **Why Use Ingress Instead of a Load Balancer?**

While a load balancer can manage external traffic to services, Ingress offers more advanced features:

1. **Fine-grained access control:** Ingress supports host-based and path-based routing, allowing you to route traffic based on the requested URL or hostname.
2. **SSL termination:** Ingress can handle secure HTTPS communication, which is not inherently managed by a standard load balancer.
3. **Single entry point:** Ingress consolidates multiple services under a single IP address, whereas a load balancer typically requires one IP address per service.

A load balancer is useful for managing traffic at a higher level but lacks the fine-grained routing capabilities of Ingress.

### **Example Scenario**

Imagine you have multiple services running in your Kubernetes cluster, with each serving a different application:


- `example.com/app1` for one application
- `example.com/app2` for another

Using Ingress, you can configure routing rules to direct traffic to the correct service based on the requested path or hostname. In contrast, a load balancer can only route traffic based on ports and cannot handle URL-based routing.

## Types of Routing in Ingress


### 1. Path-based routing:

Directs traffic to different services based on the URL path.



### 2. Host-based routing:

Directs traffic to different services based on the hostname or subdomain.



## Steps to Install and Use Ingress in Kubernetes :

### 1. Install the NGINX Ingress Controller

To install the NGINX Ingress Controller, use the following command:

```
kubectl create -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.2.1/deploy/static/provider/cloud/deploy.yaml
```

### 2. Deploy the Applications

You need to deploy two applications for testing Ingress. The application manifests can be found in the following GitHub repository:

**GitHub URL:**

<https://github.com/mustafaprojectsindevops/kubernetes/tree/master/ingress>

### 3. Apply the Manifests

Execute all the files from the repository to deploy the necessary resources.

Once the applications are deployed, use the following command to verify the Ingress resource:

```
kubectl get ing
```

### 4. Retrieve the Load Balancer DNS

After approximately 30 seconds, a load balancer DNS will be provided by the ingress controller. Use this DNS to access the deployed applications.

### 5. Access the Applications

Access the applications using the DNS provided by the ingress resource, appending the appropriate path for each application:

- **DNS/nginx:** Access the NGINX application.
- **DNS/httpd:** Access the HTTPD application.

The traffic will be routed to the respective applications based on the routing rules defined in the Ingress resource.

## Example YAML File (ONE.YML)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: one
spec:
  replicas: 2
  selector:
    matchLabels:
      app: swiggy
  template:
    metadata:
      labels:
        app: swiggy
    spec:
      containers:
      - name: cont-1
        image: nginx
        ports:
        - containerPort: 80
        env:
        - name: TITLE
          value: "NGINX APP1"
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  type: NodePort
  ports:
  - port: 80
  selector:
    app: swiggy
```

```
[root@ip-172-31-89-240 ~]# git clone https://github.com/mustafaprojectsindevops/kubernetes.git
Cloning into 'kubernetes'...
remote: Enumerating objects: 22, done.
remote: Counting objects: 100% (22/22), done.
remote: Compressing objects: 100% (19/19), done.
remote: Total 22 (delta 8), reused 3 (delta 1), pack-reused 0 (from 0)
Receiving objects: 100% (22/22), 6.27 KiB | 6.27 MiB/s, done.
Resolving deltas: 100% (8/8), done.
[root@ip-172-31-89-240 ~]# ll
total 65940
drwxr-xr-x 3 root root      78 Dec 13 19:24 aws
-rw-r--r-- 1 root root 67517785 Dec 16 05:08 awscliv2.zip
drwxr-xr-x 4 root root      33 Dec 16 07:02 kubernetes
-rw-r--r-- 1 root root    401 Dec 16 06:27 pod.yml
[root@ip-172-31-89-240 ~]# cd kubernetes
[root@ip-172-31-89-240 kubernetes]# ll
total 0
drwxr-xr-x 2 root root 55 Dec 16 07:02 ingress
[root@ip-172-31-89-240 kubernetes]# cd ingress
[root@ip-172-31-89-240 ingress]# vim one.yml
[root@ip-172-31-89-240 ingress]# vim two.yml
[root@ip-172-31-89-240 ingress]# vim ingress.yml
[root@ip-172-31-89-240 ingress]# kubectl create -f ingress.yml
ingress.networking.k8s.io/k8s-ingress created
[root@ip-172-31-89-240 ingress]# kubectl create -f one.yml
```

### Example YAML File (Two.YML)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: two
spec:
  replicas: 2
  selector:
    matchLabels:
      app: zomato
  template:
    metadata:
      labels:
        app: zomato
    spec:
      containers:
        - name: cont-2
          image: httpd
          ports:
            - containerPort: 80
          env:
            - name: TITLE
              value: "APACHE APP2"
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: httpd
spec:
  type: NodePort
  ports:
    - port: 80
  selector:
    app: zomato
```

## Creating one.yml, two.yml and ingress.yml

```
[root@ip-172-31-89-240 ingress]# vim one.yml
[root@ip-172-31-89-240 ingress]# vim two.yml
[root@ip-172-31-89-240 ingress]# vim ingress.yml
[root@ip-172-31-89-240 ingress]# kubectl create -f ingress.yml
ingress.networking.k8s.io/k8s-ingress created
[root@ip-172-31-89-240 ingress]# kubectl create -f one.yml
deployment.apps/one created
service/nginx created
[root@ip-172-31-89-240 ingress]# kubectl create -f two.yml
deployment.apps/two created
service/httpd created
[root@ip-172-31-89-240 ingress]# kubectl get ingress
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
k8s-ingress	nginx	*	a39935ccfd75f44c4a94f725c308c0c7-1681854524.us-east-1.elb.amazonaws.com	80	84s

```
[root@ip-172-31-89-240 ingress]#
```

Here, we are accessing the NGINX server interface through the LoadBalancer's DNS.

a39935ccfd75f44c4a94f725c308c0c7-1681854524.us-east-1.elb.amazonaws.com

## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](https://nginx.org).  
Commercial support is available at [nginx.com](https://nginx.com).

*Thank you for using nginx.*

Here, we are accessing the applications using the LoadBalancer, but the key difference from the previous setup is that now we explicitly add the path /httpd to the end of the LoadBalancer URL. This ensures it directly opens the HTTPD page. Previously, when we used only the LoadBalancer URL without a specific path, it would default to opening the NGINX default page.

← → ↻ ⚠ Not secure a39935ccfd75f44c4a94f725c308c0c7-1681854524.us-east-1.elb.amazonaws.com/httpd

## It works!

Example of an Ingress YAML file (ingress.yml):

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: k8s-ingress
  annotations:
    nginx.ingress.kubernetes.io/ssl-redirect: "false"
    nginx.ingress.kubernetes.io/use-regex: "true"
    nginx.ingress.kubernetes.io/rewrite-target: /$2
spec:
  ingressClassName: nginx
  rules:
    - http:
        paths:
          - path: /nginx(/|$) (.*)
            pathType: Prefix
            backend:
              service:
                name: nginx
                port:
                  number: 80
          - path: /httpd(/|$) (.*)
            pathType: Prefix
            backend:
              service:
```

```
                name: httpd
                port:
                  number: 80
    - path: /(.*)
      pathType: Prefix
      backend:
        service:
          name: nginx
          port:
            number: 80
```